

TRAINTHEML MODEL ONIBM

TeamID	PNT2022TMID30453
ProjectName	CarResalevaluePrediction

TRAINTHEML MODELONIBM

IBM Watson Studio

Search in your workspaces

Buy

4

SUGANTH V's Account

Dallas

SV

Deployments /

CAR

OverviewAssetsDeploymentsJobsManage

🔍 Search

Name	Type	Status	Asset	Last modified	
Model_IBM	Online	Deployed	resale_model	2 hours ago SUGANTH V (You)	

Items per page: 201-1 of 1 items1 of 1 pages

```
import pandas as
pdimport numpy as
npimportmatplotlibbaspl
t
from sklearn.preprocessing import
LabelEncoderimport pickle
print("IMPORTEDREQUIREDLIBRARIES")
#df=pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv",header=0,sep=',',
,encoding='Latin1',low_memory=False)
#df.head()
import os,
typesimportpandasasp
d
from botocore.client import
Configimport ibm_boto3
import io
def iter(self): return0

#@hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your
credentials.#You mightwanttoremove thosecredentials before you sharethe notebook.
cos_client=ibm_boto3.client(service_name='s3',
ibm_api_key_id='DT15l-
```

IL0017uhnUGwXyhG_Eort5gohoW6XJTNoT3RKk',ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",

```
config=Config(signature_version='oauth'),endpoint_url='https://s3.private.
us.cloud-object-storage.appdomain.cloud')
```

```
bucket = 'carresalevalueprediction-donotdelete-pr-
yuhtzmzidi0ka1p'object_key= 'autos.csv'
```

```
body=cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv((io.BytesIO(body['Body'].read())) , header=0 , sep=',',
,encoding='Latin1',low_memory=False)df.head()
#df=pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv",header=0,sep=',',
,encoding='Latin1',low_memory=False)
#df.head()
import os,
typesimportpandasasp
d
from botocore.client import
Configimport ibm_boto3
import io
def iter(self): return0
```

```
#@hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your
credentials.#You mightwanttoremove thosecredentials before you sharethe notebook.
cos_client=ibm_boto3.client(service_name='s3',
ibm_api_key_id='DT151-
IL0017uhnUGwXyhG_Eort5gohoW6XJTNoT3RkK',ibm_auth_endpoint="http
s://iam.cloud.ibm.com/oidc/token",config=Config(signature_version='oauth'),e
ndpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')
```

```
bucket = 'carresalevalueprediction-donotdelete-pr-
yuhtzmzidi0ka1p'object_key= 'autos.csv'
```

```
body=cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv((io.BytesIO(body['Body'].read())) , header=0 , sep=',',
,encoding='Latin1',low_memory=False)df.head()
print(df.seller.value_counts())
df[df.seller
!='gewerblich']df=df.drop('sel
ler',axis=1)
```

```
print(df.offerType.value_counts())df[df.
offerType
!='Gesuch']df=df.drop('offerType',axis=
1)print(df.shape)
df=df[(df.powerPS>50) &
(df.powerPS<900)]print(df.shape)
df=df[(df.yearOfRegistration>=1950)&(df.yearOfRegistration<2022)]print(d
f.shape)
df.drop(['name','abtest','dateCrawled','nrOfPictures','lastSeen','postalCode','dateCreated'],axis='columns',inplace=True)ne
w_df=df.copy()new_df=new_df.drop_duplicates(['price','vehicleType','yearOfRegistration','gearbox','powerPS','model',
kilometer','monthOfRegistration','fuelType','notRepairedDamage'])new_df.gearbox.replace(('manuell','automatik'),('ma
nual','automatic'),inplace=True)new_df.fuelType.replace(('benzin','andere','elektro'),('petrol','others','electric'),inplace=T
rue)new_df.vehicleType.replace(('kleinwagen','cabrio','kombi','andere'),('samllcar','convertible','combination','others'),in
place=True)
```

```
new_df.notRepairedDamage.replace(('ja','nein'),('Yes','No'),inplace=True)new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]
```

```
new_df['notRepairedDamage'].fillna(value='not-declared',inplace=True)new_df['fuelType'].fillna(value='not-declared',inplace=True)new_df['gearbox'].fillna(value='not-declared',inplace=True)new_df['vehicleType'].fillna(value='not-declared',inplace=True)new_df['model'].fillna(value='not-declared',inplace=True)
from ibm_watson_machine_learning import APIClientwml_credentials={
    "url":"https://us-south.ml.cloud.ibm.com",
    "apikey":"hEAn_mcoP3u_-ZjagjeqlxDayqUiETpYVYWdR1OLKAby"
}
client=APIClient(wml_credentials)
def guide_from_space_name(client,
    space_name):space=client.spaces.get_details()
#    print(space)
    return(next(item for item in space['resources'] if
item['entity']['name']==space_name)['metadata']['id'])space_uid=guide_from_space_name(client,'CAR')
print("Space UID"+
space_uid)client.set.default_space(
space_uid)client.software_specifications.list()
software_spec_uid= client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")software_spec_uid
print(new_df)labels=['gearbox','notRepairedDamage','model','brand','fuelType','vehicleType']
```

```
mapper={ }foriinlabels:
    mapper[i]=LabelEncoder()mapper[i].fit(new_df[i])
    tr=mapper[i].transform(new_df[i])np.save(str('classes'+i+'.npy'),mapper[i].classes_)print(i,":",mapper[i])
    new_df.loc[:,i+'_labels']=pd.Series(tr,index=new_df.index)
```

```
labeled=new_df[['price','yearOfRegistration','powerPS','kilometer','monthOfRegistration']+[x+'_'+labels"forxinlabels]]
print(labeled.columns)Y=labeled.iloc[:,0].valuesX=labeled.iloc[:,1:].values
```

```
Y=Y.reshape(-1,1)
fromsklearn.model_selectionimportcross_val_score,train_test_split
X_train , X_test, Y_train , Y_test =
train_test_split(X,Y,test_size=0.3,random_state=3)from
sklearn.ensembleimportRandomForestRegressor
fromsklearn.metricsimportr2_score
regressor=RandomForestRegressor(n_estimators=1000,max_depth=10,random_state=34)

regressor.fit(X_train,np.ravel(Y_train,order='C'))y_pred =
regressor.predict(X_test)print(r2_score(Y_test,y_pred))
filename='resale_model.sav'pickle.dump(regressor,open(filename,'wb'))
```

```

model_details =
    client.repository.store_model(model=regressor,meta_props={ client.repository.Mo
delMetaNames.NAME:
    "resale_model",client.repository.ModelMetaNames.SOFTWARE_SPEC_UID:
    software_spec_uid,client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id =
client.repository.get_model_id(model_details)model_id
X_train[0]
regressor.predict([[2012.0,179.0,'1500000',12.0,0, 0,30,1,1,4]])

```

Projects / CAR RESALE VALUE PREDICTION / Model Building

```

In [3]: import pandas as pd
import numpy as np
import matplotlib as plt
from sklearn.preprocessing import LabelEncoder
import pickle
print("IMPORTED REQUIRED LIBRARIES")

```

IMPORTED REQUIRED LIBRARIES

```

In [4]: # df = pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv", header=0 , sep=',', encoding='Latin1',low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='DT151-1L0017uHnUGwXyhG_Eort5gohow6XJTN0T3Rkk',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv(io.BytesIO(body['Body'].read()), header=0 , sep=',', encoding='Latin1',low_memory=False)
df.head()

```

Out[4]:	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand
0	24-03-2016 11.52	Golf_3_1_6	privat	Angebot	480.0	test	NaN	1993.0	manuell	0.0	golf	150000	0.0	benzin	volkswagen

Projects / CAR RESALE VALUE PREDICTION / Model Building

Out[4]:	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand
0	24-03-2016 11.52	Golf_3_1_6	privat	Angebot	480.0	test	NaN	1993.0	manuell	0.0	golf	150000	0.0	benzin	volkswagen
1	24-03-2016 10.58	A5_Sportback_2_7_Tdi	privat	Angebot	18300.0	test	coupe	2011.0	manuell	190.0	NaN	125000	5.0	diesel	audi
2	14-03-2016 12.52	Jeep_Grand_Cherokee_Overland	privat	Angebot	9800.0	test	suv	2004.0	automatik	163.0	grand	125000	8.0	diesel	jeep
3	17-03-2016 16.54	GOLF_4_1_4_3TÜRER	privat	Angebot	1500.0	test	kleinwagen	2001.0	manuell	75.0	golf	150000	6.0	benzin	volkswagen
4	31-03-2016 17.25	Skoda_Fabia_1_4_TDI_PD_Classic	privat	Angebot	3600.0	test	kleinwagen	2008.0	manuell	69.0	fabia	90000	7.0	diesel	skoda

```

In [5]: print(df.seller.value_counts())
df[df.seller != 'gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType != 'Gesuch']
df=df.drop('offerType',axis=1)

privat      371534
gewerblich    3
golf         1
Name: seller, dtype: int64
Angebot     371525
Gesuch       12
150000       1
Name: offerType, dtype: int64

```

```
In [6]: print(df.shape)
```

```

In [5]: print(df.seller.value_counts())
df[df.seller != 'gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType != 'Gesuch']
df=df.drop('offerType',axis=1)

privat      371534
gewerblich    3
golf         1
Name: seller, dtype: int64
Angebot      371525
Gesuch       12
150000       1
Name: offerType, dtype: int64

In [6]: print(df.shape)
df=df[(df.powerPS>50) & (df.powerPS<900)]
print(df.shape)
df=df[(df.yearOfRegistration>=1950)&(df.yearOfRegistration<2022)]
print(df.shape)

(371539, 18)
(319717, 18)
(319649, 18)

In [7]: df.drop(['name','abtest','dateCrawled','nrOfPictures','lastSeen','postalCode','dateCreated'], axis='columns',inplace=True)

In [8]: new_df=df.copy()
new_df=new_df.drop_duplicates(['price','vehicleType','yearOfRegistration','gearbox','powerPS','model','kilometer','monthOfRegistration','fuelType','notRepairedDamage'])

In [9]: new_df.gearbox.replace(('manuell','automatik'),('manual','automatic'),inplace=True)
new_df.fuelType.replace(('benzin','andere','elektro'),('petrol','others','electric'),inplace=True)
new_df.vehicleType.replace(('kleinwagen','cabrio','kombi','andere'),('small car','convertible','combination','others'),inplace=True)
new_df.notRepairedDamage.replace(('ja','nein'),('Yes','No'),inplace=True)

In [10]: new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

```

```

In [7]: df.drop(['name','abtest','dateCrawled','nrOfPictures','lastSeen','postalCode','dateCreated'], axis='columns',inplace=True)

In [8]: new_df=df.copy()
new_df=new_df.drop_duplicates(['price','vehicleType','yearOfRegistration','gearbox','powerPS','model','kilometer','monthOfRegistration','fuelType','notRepairedDamage'])

In [9]: new_df.gearbox.replace(('manuell','automatik'),('manual','automatic'),inplace=True)
new_df.fuelType.replace(('benzin','andere','elektro'),('petrol','others','electric'),inplace=True)
new_df.vehicleType.replace(('kleinwagen','cabrio','kombi','andere'),('small car','convertible','combination','others'),inplace=True)
new_df.notRepairedDamage.replace(('ja','nein'),('Yes','No'),inplace=True)

In [10]: new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

new_df['notRepairedDamage'].fillna(value='not-declared',inplace=True)
new_df['fuelType'].fillna(value='not-declared',inplace=True)
new_df['gearbox'].fillna(value='not-declared',inplace=True)
new_df['vehicleType'].fillna(value='not-declared',inplace=True)
new_df['model'].fillna(value='not-declared',inplace=True)

In [11]: from ibm_watson_machine_learning import APIClient
wml_credentials={
    "url":"https://us-south.ml.cloud.ibm.com",
    "apikey":"hEAn_mcoP3u_-ZjagjeqlxDayQUIETpYVYndR1OLKAbY"
}
client =APIClient(wml_credentials)

In [12]: def guide_from_space_name(client, space_name):
    space = client.spaces.get_details()
    # print(space)
    return(next(item for item in space['resources'] if item['entity']['name']==space_name)['metadata']['id'])

In [13]: space_uid=guide_from_space_name(client,'CAR')
print("Space UID"+ space_uid)

Space UIDbe467bbb-03a2-40e7-bf5a-91836e346951

In [14]: client.set.default_space(space_uid)

Out[14]: 'SUCCESS'

```

```
In [15]: client.software_specifications.list()

-----
NAME                               ASSET_ID                               TYPE
default_py3.6                      0062b8c9-8b7d-44a0-a9b9-46c416adcbd9 base
kernel-spark3.2-scala2.12          020d69ce-7ac1-5e68-ac1a-31189867356a base
pytorch-onnx_1.3-py3.7-edt         069ea134-3346-5748-b513-49120e15d288 base
scikit-learn_0.20-py3.6            09c5a1d0-9c1e-4473-a344-eb7b665fff687 base
spark-mllib_3.0-scala_2.12         09f4cff0-90a7-5899-b9ed-1ef348aebdee base
pytorch-onnx_rt22.1-py3.9          0b848dd4-e681-5599-be41-b5f6fccc6471 base
ai-function_0.1-py3.6              0cdb0f1e-5376-4f4d-92dd-da3b69aa9bda base
shiny-r3.6                         0e6e79df-875e-4f24-8ae9-62dcc2148306 base
tensorflow_2.4-py3.7-horovod       1092590a-307d-563d-9b62-4eb7d64b3f22 base
pytorch_1.1-py3.6                  10ac12d6-6b30-4ccd-8392-3e922c096a92 base
tensorflow_1.15-py3.6-ddl          111e41b3-de2d-5422-a4d6-bf776828c4b7 base
autoai-kb_rt22.2-py3.10            125b6d9a-5b1f-5e8d-972a-b251688ccf40 base
runtime-22.1-py3.9                 12b83a17-24d8-5082-900f-0ab31fbfd3cb base
scikit-learn_0.22-py3.6            154010fa-5b3b-4ac1-82af-4d5ee5abbc85 base
default_r3.6                       1b70aec3-ab34-4b87-8aa0-a4a3c8296a36 base
pytorch-onnx_1.3-py3.6             1bc6029a-cc97-56da-b8e0-39c3880dbbe7 base
kernel-spark3.3-r3.6               1c9e5454-f216-59dd-a20e-474a5cdf5988 base
pytorch-onnx_rt22.1-py3.9-edt      1d362186-7ad5-5b59-8b6c-9d0880bde37f base
tensorflow_2.1-py3.6               1eb25b84-d6ed-5dde-b6a5-3fbd1665666 base
spark-mllib_3.2                     20047f72-0a98-58c7-9ff5-a77b012eb8f5 base
tensorflow_2.4-py3.8-horovod       217c16f6-178f-56bf-824a-b19f20564c49 base
runtime-22.1-py3.9-cuda            26215f05-08c3-5a41-a1b0-da66306ce658 base
do_py3.8                           295addb5-9ef9-547e-9bf4-92ae3563e720 base
autoai-ts_3.8-py3.8                2aa0c932-798f-5ae9-abd6-15e0c2402fb5 base
tensorflow_1.15-py3.6              2b73a275-7cbf-420b-a912-eae7f436e0bc base
kernel-spark3.3-py3.9              2b7961e2-e3b1-5a8c-a491-482c8368839a base
pytorch_1.2-py3.6                  2c8ef57d-2687-4b7d-acce-01f94976dac1 base
spark-mllib_2.3                     2e51f700-bca0-4b0d-88dc-5c6791338875 base
pytorch-onnx_1.1-py3.6-edt         32983cea-3f32-4400-8965-dde8748d67e base
spark-mllib_3.0-py37                36507ebe-8770-5bba-ab2a-eafe787600e9 base
spark-mllib_2.4                     390d21f8-e58b-4fac-9c55-d7ceda621326 base
autoai-ts_rt22.2-py3.10            396b2e83-0953-5b86-9a55-7ce1628a406f base
xgboost_0.82-py3.6                 39e31acd-5f30-41dc-ae44-60233c80306e base
pytorch-onnx_1.2-py3.6-edt         40589d0e-7019-4e28-8daa-fb03b6f4fe12 base
pytorch-onnx_rt22.2-py3.10         40e73f55-783a-5535-b3fa-0c8b94291431 base
default_r36py38                    41c247d3-45f8-5a71-b065-8580229facf0 base
```

```
In [16]: software_spec_uid = client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")
software_spec_uid
```

Out[16]: '12b83a17-24d8-5082-900f-0ab31fbfd3cb'

```
In [17]: print(new_df)

   price  vehicleType  yearOfRegistration  gearbox  powerPS  \
1  18300.0         coupe                2011.0    manual    190.0
2   9800.0         suv                  2004.0  automatic    163.0
3   1500.0  samll car                  2001.0    manual     75.0
4   3600.0  samll car                  2008.0    manual     69.0
5    650.0  limousine                  1995.0    manual    102.0
...     ...         ...                  ...     ...     ...
371531  3200.0  limousine                  2004.0    manual    225.0
371535  1199.0  convertible                2000.0  automatic    101.0
371536  9200.0         bus                  1996.0    manual    102.0
371537  3400.0  combination                2002.0    manual    100.0
371538  28990.0  limousine                  2013.0    manual    320.0

   model kilometer  monthOfRegistration  fuelType  brand  \
1  not-declared   125000                5.0   diesel   audi
2      grand     125000                8.0   diesel   jeep
3      golf     150000                6.0   petrol  volkswagen
4      fabia     90000                7.0   diesel   skoda
5        3er     150000               10.0   petrol   bmw
...     ...         ...                  ...     ...     ...
371531      leon     150000                5.0   petrol   seat
371535    fortwo     125000                3.0   petrol   smart
371536  transporter     150000                3.0   diesel  volkswagen
371537      golf     150000                6.0   diesel  volkswagen
371538    m_reihe     50000                8.0   petrol   bmw

   notRepairedDamage
1                Yes
2      not-declared
3                  No
4                  No
5                Yes
```

```
In [21]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
regressor = RandomForestRegressor(n_estimators = 1000, max_depth = 10, random_state = 34)
```

```
In [27]: regressor.predict([[2012.0, 179.0, '1500000', 12.0, 0, 0, 30, 1, 1, 4]])
```



```
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)
regressor.fit(X_train, np.ravel(Y_train,order='C'))
```

Out[21]: RandomForestRegressor(max_depth=10, n_estimators=1000, random_state=34)

```
In [22]: y_pred = regressor.predict(X_test)
print(r2_score(Y_test,y_pred))
0.8310350387286918
```

```
In [23]: filename='resale_model.sav'
pickle.dump(regressor,open(filename,'wb'))
```

```
In [24]: model_details = client.repository.store_model(model=regressor,meta_props={
    client.repository.ModelMetaNames.NAME: "resale_model",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid,
    client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id = client.repository.get_model_id(model_details)
```

In [25]: model_id

Out[25]: 'cd8479e0-66e4-454e-aece-4824fe9d71bd'

In [26]: X_train[0:]

```
Out[26]: array([[2005.0, 179.0, '150000', ..., 1, 1, 4],
 [1997.0, 60.0, '150000', ..., 38, 7, 4],
 [2003.0, 170.0, '150000', ..., 2, 7, 1],
 ...,
 [2009.0, 174.0, '125000', ..., 25, 7, 7],
 [2000.0, 136.0, '150000', ..., 20, 7, 1],
 [2013.0, 170.0, '40000', ..., 1, 7, 8]], dtype=object)
```

```
In [27]: regressor.predict([[2012.0, 179.0, '150000', 12.0, 0, 0, 30, 1, 1, 4]])
```

Out[27]: array([18518.37919135])